

Proximal Gradient Method. Proximal operator

Seminar

Optimization for ML. Faculty of Computer Science. HSE University

Regularized / Composite Objectives

Many nonsmooth problems take the form


$$\min_{x \in \mathbb{R}^n} \varphi(x) = f(x) + r(x)$$

- **Lasso, L1-LS, compressed sensing**

$$f(x) = \frac{1}{2} \|Ax - b\|_2^2, r(x) = \lambda \|x\|_1$$

- **L1-Logistic regression, sparse LR**

$$f(x) = -y \log h(x) - (1-y) \log(1-h(x)), r(x) = \lambda \|x\|_1$$



Non-smooth convex optimization lower bounds

convex (non-smooth)

$$f(x_k) - f^* \sim \mathcal{O}\left(\frac{1}{\sqrt{k}}\right)$$
$$k_\varepsilon \sim \mathcal{O}\left(\frac{1}{\varepsilon^2}\right)$$

strongly convex (non-smooth)

$$f(x_k) - f^* \sim \mathcal{O}\left(\frac{1}{k}\right)$$
$$k_\varepsilon \sim \mathcal{O}\left(\frac{1}{\varepsilon}\right)$$

Non-smooth convex optimization lower bounds

convex (non-smooth)	strongly convex (non-smooth)
$f(x_k) - f^* \sim \mathcal{O}\left(\frac{1}{\sqrt{k}}\right)$ $k_\varepsilon \sim \mathcal{O}\left(\frac{1}{\varepsilon^2}\right)$	$f(x_k) - f^* \sim \mathcal{O}\left(\frac{1}{k}\right)$ $k_\varepsilon \sim \mathcal{O}\left(\frac{1}{\varepsilon}\right)$

- Subgradient method is optimal for the problems above.
- One can use Mirror Descent (a generalization of the subgradient method to a possibly non-Euclidian distance) with the same convergence rate to better fit the geometry of the problem.
- However, we can achieve standard gradient descent rate $\mathcal{O}\left(\frac{1}{k}\right)$ (and even accelerated version $\mathcal{O}\left(\frac{1}{k^2}\right)$) if we will exploit the structure of the problem.

Proximal operator

i Proximal operator

For a convex set $E \in \mathbb{R}^n$ and a convex function $f : E \rightarrow \mathbb{R}$ operator $\text{prox}_f(x)$ s.t.

$$\text{prox}_f(x) = \underset{y \in E}{\operatorname{argmin}} \left[f(y) + \frac{1}{2} \|y - x\|_2^2 \right]$$

is called **proximal operator** for function f at point x

From projections to proximity

Let $\mathbb{I}_S$ be the indicator function for closed, convex S . Recall orthogonal projection $\pi_S(y)$

From projections to proximity

Let $\mathbb{I}_S$ be the indicator function for closed, convex S . Recall orthogonal projection $\pi_S(y)$

$$\pi_S(y) := \arg \min_{x \in S} \frac{1}{2} \|x - y\|_2^2.$$

From projections to proximity

Let $\mathbb{I}_S$ be the indicator function for closed, convex S . Recall orthogonal projection $\pi_S(y)$

$$\pi_S(y) := \arg \min_{x \in S} \frac{1}{2} \|x - y\|_2^2.$$

With the following notation of indicator function

$$\mathbb{I}_S(x) = \begin{cases} 0, & x \in S, \\ \infty, & x \notin S, \end{cases}$$

Rewrite orthogonal projection $\pi_S(y)$ as

$$\pi_S(y) := \arg \min_{x \in \mathbb{R}^n} \frac{1}{2} \|x - y\|^2 + \mathbb{I}_S(x).$$

From projections to proximity

Let $\mathbb{I}_S$ be the indicator function for closed, convex S . Recall orthogonal projection $\pi_S(y)$

$$\pi_S(y) := \arg \min_{x \in S} \frac{1}{2} \|x - y\|_2^2.$$

With the following notation of indicator function

$$\mathbb{I}_S(x) = \begin{cases} 0, & x \in S, \\ \infty, & x \notin S, \end{cases}$$

Rewrite orthogonal projection $\pi_S(y)$ as

$$\pi_S(y) := \arg \min_{x \in \mathbb{R}^n} \frac{1}{2} \|x - y\|^2 + \mathbb{I}_S(x).$$

Proximity: Replace $\mathbb{I}_S$ by some convex function!

$$\text{prox}_r(y) = \text{prox}_{r,1}(y) := \arg \min \frac{1}{2} \|x - y\|^2 + r(x)$$

Proximal Gradient Method

💡 Proximal Gradient Method Theorem

Consider the proximal gradient method

$$x_{k+1} = \text{prox}_{\alpha r}(x_k - \alpha \nabla f(x_k))$$

for the criterion $\phi(x) = f(x) + r(x)$ s.t.:

- f is convex, differentiable with Lipschitz gradients;
- r is convex and prox-friendly.

Then Proximal Gradient Method with fixed step size $\alpha = \frac{1}{L}$ converges with rate $O(\frac{1}{k})$

Examples of Proximal Mappings

Quadratic Function ($A \geq 0$)

$$f(x) = \frac{1}{2}x^T Ax + b^T x + c,$$

Examples of Proximal Mappings

Quadratic Function ($A \geq 0$)

$$f(x) = \frac{1}{2}x^T Ax + b^T x + c,$$

$$\text{prox}_{tf}(x) = (I + tA)^{-1}(x - tb)$$

Euclidean Norm

$$f(x) = \|x\|_2,$$

Examples of Proximal Mappings

Quadratic Function ($A \geq 0$)

$$f(x) = \frac{1}{2}x^T A x + b^T x + c,$$

$$\text{prox}_{t f}(x) = (I + tA)^{-1}(x - tb)$$

Euclidean Norm

$$f(x) = \|x\|_2,$$

$$\text{prox}_{t f}(x) = \begin{cases} (1 - t/\|x\|_2)x & \text{if } \|x\|_2 \geq t \\ 0 & \text{otherwise} \end{cases}$$

Logarithmic Barrier

$$f(x) = -\sum_{i=1}^n \log x_i,$$

Examples of Proximal Mappings

Quadratic Function ($A \geq 0$)

$$f(x) = \frac{1}{2}x^T Ax + b^T x + c,$$

$$\text{prox}_{tf}(x) = (I + tA)^{-1}(x - tb)$$

Euclidean Norm

$$f(x) = \|x\|_2,$$

$$\text{prox}_{tf}(x) = \begin{cases} (1 - t/\|x\|_2)x & \text{if } \|x\|_2 \geq t \\ 0 & \text{otherwise} \end{cases}$$

Logarithmic Barrier

$$f(x) = -\sum_{i=1}^n \log x_i,$$

$$\text{prox}_{tf}(x)_i = \frac{x_i + \sqrt{x_i^2 + 4t}}{2}, \quad i = 1, \dots, n$$

Simple Calculus Rules for Proximal Mappings

Separable Sum

$$f\left(\begin{bmatrix} x \\ y \end{bmatrix}\right) = g(x) + h(y), \quad \text{prox}_f\left(\begin{bmatrix} x \\ y \end{bmatrix}\right) = \begin{bmatrix} \text{prox}_g(x) \\ \text{prox}_h(y) \end{bmatrix}$$

Scaling and Translation of Argument (for $a \neq 0$)

$$f(x) = g(ax + b), \quad \text{prox}_f(x) = \frac{1}{a} \left(\text{prox}_{a^2g}(ax + b) - b \right)$$

Right Scalar Multiplication (with $\lambda > 0$)

$$f(x) = \lambda g(x/\lambda), \quad \text{prox}_f(x) = \lambda \text{prox}_{\lambda^{-1}g}(x/\lambda)$$

ISTA and FISTA


Methods for solving problems involving $L1$ regularization (e.g. Lasso).

ISTA (Iterative Shrinkage-Thresholding Algorithm)

- Step:

$$x_{k+1} = \text{prox}_{\alpha\lambda||\cdot||_1}(x_k - \alpha\nabla f(x_k))$$

- Convergence: $O(\frac{1}{k})$



Composite optimization

FISTA (Fast Iterative Shrinkage-Thresholding Algorithm)

- Step:

$$x_{k+1} = \text{prox}_{\alpha\lambda||\cdot||_1}(y_k - \alpha\nabla f(y_k)),$$

$$t_{k+1} = \frac{1 + \sqrt{1 + 4t_k^2}}{2},$$

$$y_{k+1} = x_{k+1} + \frac{t_k - 1}{t_{k+1}}(x_{k+1} - x_k)$$

- Convergence: $O(\frac{1}{k^2})$

Problem 1. ReLU in prox

i Question

Find the $\text{prox}_f(x)$ for $f(x) = \lambda \max(0, x)$:

$$\text{prox}_{\lambda \max(0, \cdot)}(x) = \underset{y \in \mathbb{R}}{\text{argmin}} \left[\frac{1}{2} \|y - x\|^2 + \lambda \max(0, y) \right]$$

Problem 2. Grouped l_1 -regularizer

i Question

Find the $\text{prox}_f(x)$ for $f(x) = \|x\|_{1/2} = \sum_{g=0}^G \|x_g\|_2$ where $x \in \mathbb{R}^n = [\underbrace{x_1, x_2, \dots}_1, \underbrace{\dots}_g, \dots, \underbrace{x_{n-2}, x_{n-1}, x_n}_G]$:

$$\text{prox}_{\|x\|_{1/2}}(x) = \underset{y \in \mathbb{R}}{\operatorname{argmin}} \left[\frac{1}{2} \|y - x\|_2^2 + \sum_{g=0}^G \|y_g\|_2 \right]$$

Linear Least Squares with L_1 -regularizer

Proximal Methods Comparison for Linear Least Squares with L_1 -regularizer  Open in Colab.

Image Denoising using ISTA

Problem: Recover clean image x_{true} from noisy $y = x_{\text{true}} + \mathbf{n}$

Key Idea Clean images have sparse wavelet coefficients $\alpha = Wx_{\text{true}}$ (mostly zeros), where W is the Wavelet Transform operator. Noise coefficients are not sparse.

Optimization Approach: Find coefficients α minimizing a composite objective:

$$\min_{\alpha} \underbrace{\frac{1}{2} \|W^T \alpha - y\|_2^2}_{\text{Data Fidelity}} + \underbrace{\lambda \|\alpha\|_1}_{\text{Sparsity Penalty}}$$

where W^T is the inverse Wavelet Transform and λ balances the terms.

Result: The denoised image is obtained from the final coefficients α^* :

$$x_{\text{denoised}} = W^T \alpha^*$$

Image Denoising using ISTA

 Open Git.

